

# Small data and deep learning

This practical session proposes to study several techniques to improve training performance in the challenging context where few data and resources are available.

In [1]:

```
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

from tqdm import tqdm
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

## Introduction

Assume we are in a context where few "gold" labeled data are available for training, say  $\mathcal{X}_{\text{train}} \triangleq \{(x_n, y_n)\}_{n \leq N_{\text{train}}}$ , where  $N_{\text{train}}$  is small. A large test set  $\mathcal{X}_{\text{test}}$  is available. A large amount of unlabeled data,  $\mathcal{X}_{\text{nolabel}}$ , is available. We also assume that we have a limited computational budget (e.g., no GPUs).

For each question, write a commented *Code* or a complete answer as a *Markdown*. When the objective of a question is to report a CNN accuracy, please use the following format to report it, at the end of the question :

Model	Number of epochs	Train accuracy	Test accuracy
XXX	XXX	XXX	XXX

If applicable, please add the field corresponding to the **Accuracy on Full Data** as well as a link to the **Reference paper** you used to report those numbers. (You do not need to train a CNN on the full CIFAR10 dataset.)

In your final report, please keep the logs of each training procedure you used. We will only run this jupyter if we have some doubts on your implementation.

**The total file sizes should be reasonable (feasible with 2MB only!). You will be asked to hand in the notebook, together with any necessary files required to run it if any.**

To run your experiments, you can use the same local installation as for previous TPs, or otherwise <https://colab.research.google.com/>.

## Training set creation

**Question 1 (2 points) :** Propose a dataloader that will only use the first 100 samples of the CIFAR-10 training set.

*Hint :* You can modify the file located at

<https://github.com/pytorch/vision/blob/master/torchvision/datasets/cifar.py> or use the information from <https://pytorch.org/vision/stable/datasets.html>

This is our dataset  $\mathcal{X}_{\text{train}}$ , it will be used until the end of this project. The remaining samples correspond to  $\mathcal{X}_{\text{nolabel}}$ . The testing set  $\mathcal{X}_{\text{test}}$  corresponds to the whole testing set of CIFAR-10.

```
In [3]: transform_train = transforms.Compose([
    transforms.ToTensor(),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
])

batch_size = 10

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform_train)
nolabelset = torch.utils.data.Subset(trainset, list(range(100, len(trainset))))
trainset = torch.utils.data.Subset(trainset, list(range(100)))
nolabel_dataloader = torch.utils.data.DataLoader(nolabelset, batch_size=batch_size, shuffle=True)
train_dataloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform_test)
test_dataloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False)

len(train_dataloader), len(nolabel_dataloader), len(test_dataloader)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

Extracting ./data/cifar-10-python.tar.gz to ./data  
Files already downloaded and verified

Out[3]: (10, 4990, 1000)

## Testing procedure

**Question 2 (1.5 points) :** Explain why the evaluation of the training procedure is difficult.

Propose several solutions.

**Answer:**

The evaluation of the training procedure is difficult since there is a risk of overfitting the training data after a few epochs of the training process. In particular, in this case we only have 100 training samples which make the risks for overfitting quite high.

A common method to evaluate the training procedure is to include a validation set in the training process on which we can evaluate the loss and/or accuracy of the model on samples that are out of the training set at each epoch. This way we can detect when the model starts to overfit. However, in this case, we only have 100 training samples so we might not be able to afford cutting a validation set in the training set. We could still use the testing set to evaluate the model at each (few) epoch(s) but considering the size of our testing set, this might be too costly to use the whole testing set.

Another solution could be to use part of the unlabeled data. We could first apply an unsupervised learning method (such as clustering) on the training data mixed with part of the unlabeled data to infer some labels on the unlabeled data. Then, during the training on the training set, we could compare the model's predictions with the inferred labels to see if they match. The problem with this method is that the unsupervised learning could be very inaccurate and mess up the evaluation process.

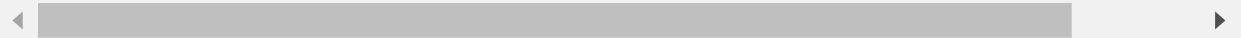
## Raw approach: the baseline

In this section, the goal is to train a CNN on  $\mathcal{X}_{\text{train}}$  and compare its performance with reported numbers from the literature. You will have to re-use and/or design a standard classification pipeline. You should optimize your pipeline to obtain the best performances (image size, data augmentation by flip, ...).

The key ingredients for training a CNN are the batch size, as well as the learning rate schedule, i.e. how to decrease the learning rate as a function of the number of epochs. A possible schedule is to start the learning rate at 0.1 and decreasing it every 30 epochs by 10. In case of divergence, reduce the learning rate. A potential batch size could be 10, yet this can be cross-validated.

You can get some baselines accuracies in this paper :

[http://openaccess.thecvf.com/content\\_cvpr\\_2018/papers/Keshari\\_Learning\\_Structure\\_and\\_CVPR\\_2018\\_paper.pdf](http://openaccess.thecvf.com/content_cvpr_2018/papers/Keshari_Learning_Structure_and_CVPR_2018_paper.pdf)  
Obviously, it is a different context for those researchers who had access to GPUs.



## ResNet architectures

**Question 3 (4 points)** : Write a classification pipeline for  $\mathcal{X}_{\text{train}}$ , train from scratch and evaluate a ResNet-18 architecture specific to CIFAR10 (details about the ResNet-18 model, originally designed for the ImageNet dataset, can be found here: <https://arxiv.org/abs/1512.03385>). Please report the accuracy obtained on the whole dataset as well as the reference paper/GitHub link.

*Hint:* You can re-use the following code: <https://github.com/kuangliu/pytorch-cifar>. During a training of 10 epochs, a batch size of 10 and a learning rate of 0.01, one obtains 40% accuracy on  $\mathcal{X}_{\text{train}}$  (~2 minutes) and 20% accuracy on  $\mathcal{X}_{\text{test}}$  (~5 minutes).

## ResNet model

In [4]:

```
class BasicBlock(nn.Module):
    expansion = 1

    def __init__(self, in_planes, planes, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = nn.Conv2d(
            in_planes, planes, kernel_size=3, stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(planes)
        self.conv2 = nn.Conv2d(planes, planes, kernel_size=3,
                            stride=1, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(planes)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_planes != self.expansion*planes:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_planes, self.expansion*planes, kernel_size=1,
                        stride=stride, bias=False),
                nn.BatchNorm2d(self.expansion*planes)
            )
```

```

        nn.Conv2d(in_planes, self.expansion*planes,
                  kernel_size=1, stride=stride, bias=False),
        nn.BatchNorm2d(self.expansion*planes)
    )

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += self.shortcut(x)
        out = F.relu(out)
        return out

class ResNet(nn.Module):
    def __init__(self, block, num_blocks, num_classes=10):
        super(ResNet, self).__init__()
        self.in_planes = 64

        self.conv1 = nn.Conv2d(3, 64, kernel_size=3,
                            stride=1, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(64)
        self.layer1 = self._make_layer(block, 64, num_blocks[0], stride=1)
        self.layer2 = self._make_layer(block, 128, num_blocks[1], stride=2)
        self.layer3 = self._make_layer(block, 256, num_blocks[2], stride=2)
        self.layer4 = self._make_layer(block, 512, num_blocks[3], stride=2)
        self.linear = nn.Linear(512*block.expansion, num_classes)

    def _make_layer(self, block, planes, num_blocks, stride):
        strides = [stride] + [1]*(num_blocks-1)
        layers = []
        for stride in strides:
            layers.append(block(self.in_planes, planes, stride))
            self.in_planes = planes * block.expansion
        return nn.Sequential(*layers)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.layer4(out)
        out = F.avg_pool2d(out, 4)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

def ResNet18():
    return ResNet(BasicBlock, [2, 2, 2, 2])

```

## Training

In [5]:

```

model = ResNet18()
model.to(device)

epochs = 100
criterion = nn.CrossEntropyLoss()
lr = 0.1
optimizer = optim.Adam(model.parameters(), lr=lr)
lr_update_period = 30
milestones = list(range(0, epochs, lr_update_period))[1:]
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones, gamma=0.1)

```

In [6]:

```

def train(model, train_dataloader, criterion, optimizer, scheduler):
    total_train_losses = []
    total_train_accs = []

    for epoch in range(1, epochs+1):
        model.train()
        train_losses = []

        correct = 0
        for i, batch, in enumerate(tqdm(train_dataloader)):
            img_batch, lbl_batch = batch
            img_batch, lbl_batch = img_batch.to(device), lbl_batch.to(device)

            optimizer.zero_grad()
            outputs = model(img_batch)
            loss = criterion(torch.flatten(outputs, 1), lbl_batch)
            loss.backward()
            optimizer.step()

            train_losses.append(loss.item())
            correct += (torch.flatten(outputs, 1).argmax(1) == lbl_batch).float().sum()

        print('Train (epoch {}/{})\tLoss: {:.6f}\tAccuracy: {:.2f}%'.format(
            epoch, epochs, np.mean(train_losses), 100 * correct / len(trainset)))

        train_loss_mean = np.mean(train_losses)
        total_train_losses.append(train_loss_mean)
        train_acc_mean = correct / len(trainset)
        total_train_accs.append(train_acc_mean)

        scheduler.step()

    train_loss = total_train_losses[-1]
    train_acc = total_train_accs[-1]
    print("\nTrain loss: {:.6f}\tTrain accuracy: {:.2f}%".format(train_loss, train_acc))

    plt.plot(list(range(epochs+1))[1:], total_train_losses)
    plt.title('Train loss evolution')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')

    plt.show()

    plt.plot(list(range(epochs+1))[1:], total_train_accs)
    plt.title('Train accuracy evolution')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')

    plt.show()

train(model, train_dataloader, criterion, optimizer, scheduler)

```

```

100%|██████████| 10/10 [00:00<00:00, 20.57it/s]
Train (epoch 1/100)      Loss: 14.173867 Accuracy: 10.00%
100%|██████████| 10/10 [00:00<00:00, 56.91it/s]
Train (epoch 2/100)      Loss: 2.647141  Accuracy: 13.00%
100%|██████████| 10/10 [00:00<00:00, 55.65it/s]
Train (epoch 3/100)      Loss: 2.481507  Accuracy: 14.00%
100%|██████████| 10/10 [00:00<00:00, 57.98it/s]
Train (epoch 4/100)      Loss: 2.271270  Accuracy: 11.00%
100%|██████████| 10/10 [00:00<00:00, 55.14it/s]
Train (epoch 5/100)      Loss: 2.282340  Accuracy: 15.00%
100%|██████████| 10/10 [00:00<00:00, 57.99it/s]

```

Train (epoch 6/100) Loss: 2.240300 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 58.67it/s]  
Train (epoch 7/100) Loss: 2.291780 Accuracy: 15.00%  
100% [██████████] 10/10 [00:00<00:00, 56.64it/s]  
Train (epoch 8/100) Loss: 2.250262 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 54.21it/s]  
Train (epoch 9/100) Loss: 2.247396 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 56.91it/s]  
Train (epoch 10/100) Loss: 2.238914 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 56.68it/s]  
Train (epoch 11/100) Loss: 2.258385 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 56.70it/s]  
Train (epoch 12/100) Loss: 2.242633 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 57.51it/s]  
Train (epoch 13/100) Loss: 2.253733 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 57.62it/s]  
Train (epoch 14/100) Loss: 2.212862 Accuracy: 14.00%  
100% [██████████] 10/10 [00:00<00:00, 59.24it/s]  
Train (epoch 15/100) Loss: 2.176782 Accuracy: 20.00%  
100% [██████████] 10/10 [00:00<00:00, 52.55it/s]  
Train (epoch 16/100) Loss: 2.263088 Accuracy: 15.00%  
100% [██████████] 10/10 [00:00<00:00, 58.39it/s]  
Train (epoch 17/100) Loss: 2.207663 Accuracy: 13.00%  
100% [██████████] 10/10 [00:00<00:00, 55.69it/s]  
Train (epoch 18/100) Loss: 2.268725 Accuracy: 15.00%  
100% [██████████] 10/10 [00:00<00:00, 57.45it/s]  
Train (epoch 19/100) Loss: 2.230460 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 57.53it/s]  
Train (epoch 20/100) Loss: 2.214703 Accuracy: 17.00%  
100% [██████████] 10/10 [00:00<00:00, 56.94it/s]  
Train (epoch 21/100) Loss: 2.169014 Accuracy: 16.00%  
100% [██████████] 10/10 [00:00<00:00, 56.97it/s]  
Train (epoch 22/100) Loss: 2.143315 Accuracy: 23.00%  
100% [██████████] 10/10 [00:00<00:00, 58.39it/s]  
Train (epoch 23/100) Loss: 2.176824 Accuracy: 19.00%  
100% [██████████] 10/10 [00:00<00:00, 58.89it/s]  
Train (epoch 24/100) Loss: 2.136893 Accuracy: 22.00%  
100% [██████████] 10/10 [00:00<00:00, 59.27it/s]  
Train (epoch 25/100) Loss: 2.132151 Accuracy: 24.00%  
100% [██████████] 10/10 [00:00<00:00, 58.28it/s]  
Train (epoch 26/100) Loss: 2.211286 Accuracy: 23.00%  
100% [██████████] 10/10 [00:00<00:00, 58.21it/s]  
Train (epoch 27/100) Loss: 2.191934 Accuracy: 19.00%  
100% [██████████] 10/10 [00:00<00:00, 58.00it/s]  
Train (epoch 28/100) Loss: 2.167705 Accuracy: 25.00%  
100% [██████████] 10/10 [00:00<00:00, 57.61it/s]  
Train (epoch 29/100) Loss: 2.215671 Accuracy: 18.00%  
100% [██████████] 10/10 [00:00<00:00, 56.21it/s]  
Train (epoch 30/100) Loss: 2.124735 Accuracy: 22.00%  
100% [██████████] 10/10 [00:00<00:00, 56.54it/s]  
Train (epoch 31/100) Loss: 2.007487 Accuracy: 23.00%  
100% [██████████] 10/10 [00:00<00:00, 56.90it/s]  
Train (epoch 32/100) Loss: 2.009177 Accuracy: 26.00%  
100% [██████████] 10/10 [00:00<00:00, 54.17it/s]  
Train (epoch 33/100) Loss: 1.986536 Accuracy: 25.00%  
100% [██████████] 10/10 [00:00<00:00, 57.57it/s]  
Train (epoch 34/100) Loss: 1.939341 Accuracy: 31.00%  
100% [██████████] 10/10 [00:00<00:00, 56.43it/s]  
Train (epoch 35/100) Loss: 1.925496 Accuracy: 29.00%

100%|██████████| 10/10 [00:00<00:00, 58.01it/s]  
Train (epoch 36/100) Loss: 1.894661 Accuracy: 29.00%  
100%|██████████| 10/10 [00:00<00:00, 56.98it/s]  
Train (epoch 37/100) Loss: 1.871072 Accuracy: 28.00%  
100%|██████████| 10/10 [00:00<00:00, 57.67it/s]  
Train (epoch 38/100) Loss: 1.949564 Accuracy: 32.00%  
100%|██████████| 10/10 [00:00<00:00, 57.42it/s]  
Train (epoch 39/100) Loss: 1.819986 Accuracy: 34.00%  
100%|██████████| 10/10 [00:00<00:00, 56.32it/s]  
Train (epoch 40/100) Loss: 1.820623 Accuracy: 33.00%  
100%|██████████| 10/10 [00:00<00:00, 55.84it/s]  
Train (epoch 41/100) Loss: 1.787356 Accuracy: 35.00%  
100%|██████████| 10/10 [00:00<00:00, 56.54it/s]  
Train (epoch 42/100) Loss: 1.797617 Accuracy: 33.00%  
100%|██████████| 10/10 [00:00<00:00, 59.12it/s]  
Train (epoch 43/100) Loss: 1.765094 Accuracy: 35.00%  
100%|██████████| 10/10 [00:00<00:00, 57.57it/s]  
Train (epoch 44/100) Loss: 1.810713 Accuracy: 34.00%  
100%|██████████| 10/10 [00:00<00:00, 58.35it/s]  
Train (epoch 45/100) Loss: 1.743226 Accuracy: 38.00%  
100%|██████████| 10/10 [00:00<00:00, 57.13it/s]  
Train (epoch 46/100) Loss: 1.752311 Accuracy: 35.00%  
100%|██████████| 10/10 [00:00<00:00, 58.75it/s]  
Train (epoch 47/100) Loss: 1.739355 Accuracy: 36.00%  
100%|██████████| 10/10 [00:00<00:00, 59.06it/s]  
Train (epoch 48/100) Loss: 1.680105 Accuracy: 41.00%  
100%|██████████| 10/10 [00:00<00:00, 56.95it/s]  
Train (epoch 49/100) Loss: 1.715570 Accuracy: 43.00%  
100%|██████████| 10/10 [00:00<00:00, 57.78it/s]  
Train (epoch 50/100) Loss: 1.801788 Accuracy: 38.00%  
100%|██████████| 10/10 [00:00<00:00, 58.14it/s]  
Train (epoch 51/100) Loss: 1.684847 Accuracy: 46.00%  
100%|██████████| 10/10 [00:00<00:00, 58.67it/s]  
Train (epoch 52/100) Loss: 1.680698 Accuracy: 40.00%  
100%|██████████| 10/10 [00:00<00:00, 58.96it/s]  
Train (epoch 53/100) Loss: 1.563620 Accuracy: 51.00%  
100%|██████████| 10/10 [00:00<00:00, 58.79it/s]  
Train (epoch 54/100) Loss: 1.645716 Accuracy: 42.00%  
100%|██████████| 10/10 [00:00<00:00, 58.06it/s]  
Train (epoch 55/100) Loss: 1.573211 Accuracy: 44.00%  
100%|██████████| 10/10 [00:00<00:00, 57.92it/s]  
Train (epoch 56/100) Loss: 1.453446 Accuracy: 50.00%  
100%|██████████| 10/10 [00:00<00:00, 57.25it/s]  
Train (epoch 57/100) Loss: 1.588128 Accuracy: 41.00%  
100%|██████████| 10/10 [00:00<00:00, 55.07it/s]  
Train (epoch 58/100) Loss: 1.500733 Accuracy: 52.00%  
100%|██████████| 10/10 [00:00<00:00, 56.63it/s]  
Train (epoch 59/100) Loss: 1.466129 Accuracy: 56.00%  
100%|██████████| 10/10 [00:00<00:00, 55.61it/s]  
Train (epoch 60/100) Loss: 1.356649 Accuracy: 59.00%  
100%|██████████| 10/10 [00:00<00:00, 56.79it/s]  
Train (epoch 61/100) Loss: 1.281722 Accuracy: 61.00%  
100%|██████████| 10/10 [00:00<00:00, 55.28it/s]  
Train (epoch 62/100) Loss: 1.295134 Accuracy: 58.00%  
100%|██████████| 10/10 [00:00<00:00, 56.95it/s]  
Train (epoch 63/100) Loss: 1.281352 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 58.88it/s]  
Train (epoch 64/100) Loss: 1.280813 Accuracy: 60.00%  
100%|██████████| 10/10 [00:00<00:00, 58.05it/s]  
Train (epoch 65/100) Loss: 1.316476 Accuracy: 57.00%

100%|██████████| 10/10 [00:00<00:00, 57.91it/s]  
Train (epoch 66/100) Loss: 1.251877 Accuracy: 66.00%  
100%|██████████| 10/10 [00:00<00:00, 58.61it/s]  
Train (epoch 67/100) Loss: 1.182145 Accuracy: 60.00%  
100%|██████████| 10/10 [00:00<00:00, 58.48it/s]  
Train (epoch 68/100) Loss: 1.279623 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 58.17it/s]  
Train (epoch 69/100) Loss: 1.197192 Accuracy: 61.00%  
100%|██████████| 10/10 [00:00<00:00, 57.42it/s]  
Train (epoch 70/100) Loss: 1.270977 Accuracy: 54.00%  
100%|██████████| 10/10 [00:00<00:00, 56.06it/s]  
Train (epoch 71/100) Loss: 1.171638 Accuracy: 58.00%  
100%|██████████| 10/10 [00:00<00:00, 56.35it/s]  
Train (epoch 72/100) Loss: 1.220926 Accuracy: 60.00%  
100%|██████████| 10/10 [00:00<00:00, 55.60it/s]  
Train (epoch 73/100) Loss: 1.154077 Accuracy: 67.00%  
100%|██████████| 10/10 [00:00<00:00, 57.84it/s]  
Train (epoch 74/100) Loss: 1.107821 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 56.74it/s]  
Train (epoch 75/100) Loss: 1.224511 Accuracy: 54.00%  
100%|██████████| 10/10 [00:00<00:00, 58.37it/s]  
Train (epoch 76/100) Loss: 1.164805 Accuracy: 64.00%  
100%|██████████| 10/10 [00:00<00:00, 54.15it/s]  
Train (epoch 77/100) Loss: 1.136278 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.47it/s]  
Train (epoch 78/100) Loss: 1.193771 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.21it/s]  
Train (epoch 79/100) Loss: 1.120062 Accuracy: 68.00%  
100%|██████████| 10/10 [00:00<00:00, 58.44it/s]  
Train (epoch 80/100) Loss: 1.122673 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.26it/s]  
Train (epoch 81/100) Loss: 1.063897 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 56.51it/s]  
Train (epoch 82/100) Loss: 1.035164 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 57.05it/s]  
Train (epoch 83/100) Loss: 1.170233 Accuracy: 59.00%  
100%|██████████| 10/10 [00:00<00:00, 56.25it/s]  
Train (epoch 84/100) Loss: 1.001470 Accuracy: 70.00%  
100%|██████████| 10/10 [00:00<00:00, 57.82it/s]  
Train (epoch 85/100) Loss: 1.043422 Accuracy: 64.00%  
100%|██████████| 10/10 [00:00<00:00, 57.84it/s]  
Train (epoch 86/100) Loss: 1.130207 Accuracy: 66.00%  
100%|██████████| 10/10 [00:00<00:00, 58.33it/s]  
Train (epoch 87/100) Loss: 1.151870 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 58.16it/s]  
Train (epoch 88/100) Loss: 1.134730 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.68it/s]  
Train (epoch 89/100) Loss: 1.000151 Accuracy: 70.00%  
100%|██████████| 10/10 [00:00<00:00, 57.71it/s]  
Train (epoch 90/100) Loss: 0.987956 Accuracy: 72.00%  
100%|██████████| 10/10 [00:00<00:00, 58.98it/s]  
Train (epoch 91/100) Loss: 1.078234 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.54it/s]  
Train (epoch 92/100) Loss: 1.012300 Accuracy: 71.00%  
100%|██████████| 10/10 [00:00<00:00, 56.67it/s]  
Train (epoch 93/100) Loss: 0.955418 Accuracy: 72.00%  
100%|██████████| 10/10 [00:00<00:00, 57.36it/s]  
Train (epoch 94/100) Loss: 1.030773 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 56.74it/s]

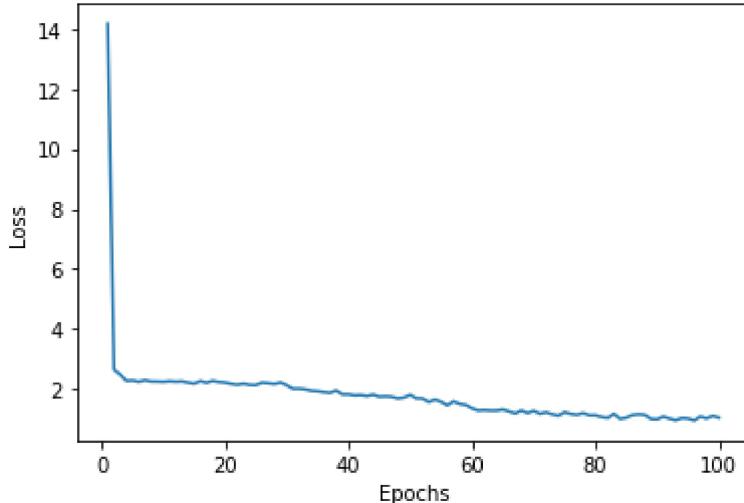
```

Train (epoch 95/100)    Loss: 1.014880  Accuracy: 70.00%
100%|██████████| 10/10 [00:00<00:00, 59.01it/s]
Train (epoch 96/100)    Loss: 0.949116  Accuracy: 72.00%
100%|██████████| 10/10 [00:00<00:00, 58.97it/s]
Train (epoch 97/100)    Loss: 1.080860  Accuracy: 68.00%
100%|██████████| 10/10 [00:00<00:00, 58.08it/s]
Train (epoch 98/100)    Loss: 1.017342  Accuracy: 64.00%
100%|██████████| 10/10 [00:00<00:00, 57.22it/s]
Train (epoch 99/100)    Loss: 1.100903  Accuracy: 67.00%
100%|██████████| 10/10 [00:00<00:00, 58.89it/s]
Train (epoch 100/100)   Loss: 1.043173  Accuracy: 64.00%

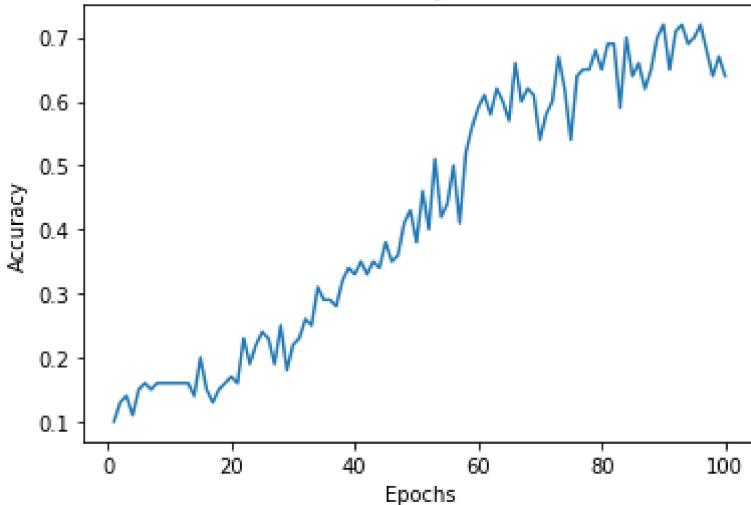
```

Train loss: 1.043173    Train accuracy: 64.00%

Train loss evolution



Train accuracy evolution



## Testing

```
In [7]: def test(model, test_dataloader):
    model.eval()
    test_losses = []
    correct = 0

    for i, batch, in enumerate(tqdm(test_dataloader)):
        img_batch, lbl_batch = batch
        img_batch, lbl_batch = img_batch.to(device), lbl_batch.to(device)

        outputs = model(img_batch)
        loss = criterion(torch.flatten(outputs, 1), lbl_batch)
```

```

        test_losses.append(loss.item())
        correct += (torch.flatten(outputs, 1).argmax(1) == lbl_batch).float().sum()

    test_loss = np.mean(test_losses)
    test_acc = correct / len(testset)

    print("\nTest loss: {:.6f}\tTest accuracy: {:.2f}%".format(test_loss, 100 * test
test(model, test_dataloader))

```

100%|██████████| 1000/1000 [00:05<00:00, 187.07it/s]  
Test loss: 2.650611      Test accuracy: 19.40%

## Results

Model	Number of epochs	Train accuracy	Test accuracy
ResNet18	100	0.64	0.1940

## Transfer learning

We propose to use pre-trained models on a classification task, in order to improve the results of our setting.

## ImageNet features

Now, we will use a model pre-trained on ImageNet and see how well it performs on CIFAR. A list of ImageNet pre-trained models is available on : <https://pytorch.org/vision/stable/models.html>

**Question 4 (3 points)** : Pick a model from the list above, adapt it for CIFAR and retrain its final layer (or a block of layers, depending on the resources to which you have access to). Report its accuracy.

## Training

In [18]:

```

model = torchvision.models.resnet18(pretrained=True)
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, 10)
model.to(device)

epochs = 100
criterion = nn.CrossEntropyLoss()
lr = 0.001
optimizer = optim.Adam(model.parameters(), lr=lr)
lr_update_period = 30
milestones = list(range(0, epochs, lr_update_period))[1:]
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones, gamma=0.1)

```

In [19]:

```
train(model, train_dataloader, criterion, optimizer, scheduler)
```

100%|██████████| 10/10 [00:00<00:00, 71.15it/s]
Train (epoch 1/100)      Loss: 2.654333      Accuracy: 24.00%
100%|██████████| 10/10 [00:00<00:00, 70.61it/s]
Train (epoch 2/100)      Loss: 2.105064      Accuracy: 42.00%
100%|██████████| 10/10 [00:00<00:00, 70.90it/s]
Train (epoch 3/100)      Loss: 1.268868      Accuracy: 59.00%

```
100%|██████████| 10/10 [00:00<00:00, 70.66it/s]
Train (epoch 4/100)    Loss: 1.155815 Accuracy: 62.00%
100%|██████████| 10/10 [00:00<00:00, 65.27it/s]
Train (epoch 5/100)    Loss: 0.929193 Accuracy: 70.00%
100%|██████████| 10/10 [00:00<00:00, 68.87it/s]
Train (epoch 6/100)    Loss: 1.029377 Accuracy: 71.00%
100%|██████████| 10/10 [00:00<00:00, 64.37it/s]
Train (epoch 7/100)    Loss: 1.046078 Accuracy: 71.00%
100%|██████████| 10/10 [00:00<00:00, 68.90it/s]
Train (epoch 8/100)    Loss: 0.789808 Accuracy: 80.00%
100%|██████████| 10/10 [00:00<00:00, 71.93it/s]
Train (epoch 9/100)    Loss: 0.842119 Accuracy: 78.00%
100%|██████████| 10/10 [00:00<00:00, 68.96it/s]
Train (epoch 10/100)   Loss: 0.585474 Accuracy: 81.00%
100%|██████████| 10/10 [00:00<00:00, 68.36it/s]
Train (epoch 11/100)   Loss: 0.875376 Accuracy: 74.00%
100%|██████████| 10/10 [00:00<00:00, 69.34it/s]
Train (epoch 12/100)   Loss: 0.931330 Accuracy: 79.00%
100%|██████████| 10/10 [00:00<00:00, 68.46it/s]
Train (epoch 13/100)   Loss: 0.679561 Accuracy: 81.00%
100%|██████████| 10/10 [00:00<00:00, 61.16it/s]
Train (epoch 14/100)   Loss: 0.594977 Accuracy: 79.00%
100%|██████████| 10/10 [00:00<00:00, 56.58it/s]
Train (epoch 15/100)   Loss: 0.762102 Accuracy: 72.00%
100%|██████████| 10/10 [00:00<00:00, 33.24it/s]
Train (epoch 16/100)   Loss: 0.564322 Accuracy: 81.00%
100%|██████████| 10/10 [00:00<00:00, 36.96it/s]
Train (epoch 17/100)   Loss: 0.594385 Accuracy: 85.00%
100%|██████████| 10/10 [00:00<00:00, 35.89it/s]
Train (epoch 18/100)   Loss: 0.390600 Accuracy: 91.00%
100%|██████████| 10/10 [00:00<00:00, 31.95it/s]
Train (epoch 19/100)   Loss: 0.535006 Accuracy: 87.00%
100%|██████████| 10/10 [00:00<00:00, 47.54it/s]
Train (epoch 20/100)   Loss: 0.530476 Accuracy: 82.00%
100%|██████████| 10/10 [00:00<00:00, 35.02it/s]
Train (epoch 21/100)   Loss: 0.317350 Accuracy: 92.00%
100%|██████████| 10/10 [00:00<00:00, 37.81it/s]
Train (epoch 22/100)   Loss: 0.351545 Accuracy: 85.00%
100%|██████████| 10/10 [00:00<00:00, 25.52it/s]
Train (epoch 23/100)   Loss: 0.546726 Accuracy: 85.00%
100%|██████████| 10/10 [00:00<00:00, 34.64it/s]
Train (epoch 24/100)   Loss: 0.373854 Accuracy: 91.00%
100%|██████████| 10/10 [00:00<00:00, 34.76it/s]
Train (epoch 25/100)   Loss: 0.681156 Accuracy: 84.00%
100%|██████████| 10/10 [00:00<00:00, 28.85it/s]
Train (epoch 26/100)   Loss: 0.406504 Accuracy: 84.00%
100%|██████████| 10/10 [00:00<00:00, 31.24it/s]
Train (epoch 27/100)   Loss: 0.368460 Accuracy: 89.00%
100%|██████████| 10/10 [00:00<00:00, 37.73it/s]
Train (epoch 28/100)   Loss: 0.271760 Accuracy: 90.00%
100%|██████████| 10/10 [00:00<00:00, 36.20it/s]
Train (epoch 29/100)   Loss: 0.119364 Accuracy: 96.00%
100%|██████████| 10/10 [00:00<00:00, 35.77it/s]
Train (epoch 30/100)   Loss: 0.333305 Accuracy: 89.00%
100%|██████████| 10/10 [00:00<00:00, 30.51it/s]
Train (epoch 31/100)   Loss: 0.403329 Accuracy: 86.00%
100%|██████████| 10/10 [00:00<00:00, 28.54it/s]
Train (epoch 32/100)   Loss: 0.090575 Accuracy: 98.00%
100%|██████████| 10/10 [00:00<00:00, 25.50it/s]
```

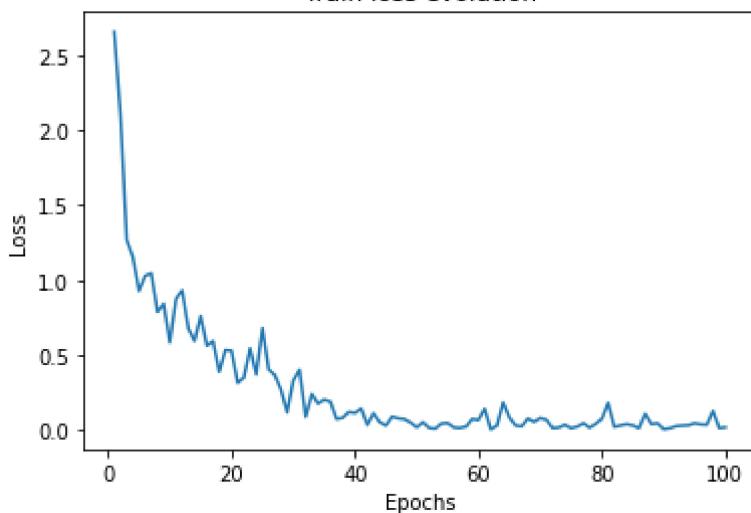
Train (epoch 33/100) Loss: 0.243140 Accuracy: 92.00%  
100% [██████████] 10/10 [00:00<00:00, 27.63it/s]  
Train (epoch 34/100) Loss: 0.178384 Accuracy: 94.00%  
100% [██████████] 10/10 [00:00<00:00, 29.38it/s]  
Train (epoch 35/100) Loss: 0.204884 Accuracy: 95.00%  
100% [██████████] 10/10 [00:00<00:00, 35.48it/s]  
Train (epoch 36/100) Loss: 0.192377 Accuracy: 93.00%  
100% [██████████] 10/10 [00:00<00:00, 34.94it/s]  
Train (epoch 37/100) Loss: 0.076209 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 38.69it/s]  
Train (epoch 38/100) Loss: 0.084062 Accuracy: 98.00%  
100% [██████████] 10/10 [00:00<00:00, 42.86it/s]  
Train (epoch 39/100) Loss: 0.122668 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 35.74it/s]  
Train (epoch 40/100) Loss: 0.115955 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 22.27it/s]  
Train (epoch 41/100) Loss: 0.146371 Accuracy: 92.00%  
100% [██████████] 10/10 [00:00<00:00, 31.08it/s]  
Train (epoch 42/100) Loss: 0.037030 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 60.58it/s]  
Train (epoch 43/100) Loss: 0.113854 Accuracy: 96.00%  
100% [██████████] 10/10 [00:00<00:00, 72.87it/s]  
Train (epoch 44/100) Loss: 0.055109 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 68.13it/s]  
Train (epoch 45/100) Loss: 0.033256 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 70.06it/s]  
Train (epoch 46/100) Loss: 0.091546 Accuracy: 98.00%  
100% [██████████] 10/10 [00:00<00:00, 67.78it/s]  
Train (epoch 47/100) Loss: 0.080521 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 65.70it/s]  
Train (epoch 48/100) Loss: 0.075182 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 70.12it/s]  
Train (epoch 49/100) Loss: 0.052392 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 70.41it/s]  
Train (epoch 50/100) Loss: 0.021225 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 68.09it/s]  
Train (epoch 51/100) Loss: 0.054994 Accuracy: 98.00%  
100% [██████████] 10/10 [00:00<00:00, 69.05it/s]  
Train (epoch 52/100) Loss: 0.016674 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 68.70it/s]  
Train (epoch 53/100) Loss: 0.011970 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 70.36it/s]  
Train (epoch 54/100) Loss: 0.045598 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 72.97it/s]  
Train (epoch 55/100) Loss: 0.049664 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 68.30it/s]  
Train (epoch 56/100) Loss: 0.019863 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 74.21it/s]  
Train (epoch 57/100) Loss: 0.016603 Accuracy: 100.00%  
100% [██████████] 10/10 [00:00<00:00, 68.24it/s]  
Train (epoch 58/100) Loss: 0.026326 Accuracy: 99.00%  
100% [██████████] 10/10 [00:00<00:00, 69.70it/s]  
Train (epoch 59/100) Loss: 0.076123 Accuracy: 97.00%  
100% [██████████] 10/10 [00:00<00:00, 69.46it/s]  
Train (epoch 60/100) Loss: 0.069347 Accuracy: 98.00%  
100% [██████████] 10/10 [00:00<00:00, 67.01it/s]  
Train (epoch 61/100) Loss: 0.144144 Accuracy: 94.00%  
100% [██████████] 10/10 [00:00<00:00, 67.35it/s]  
Train (epoch 62/100) Loss: 0.007285 Accuracy: 100.00%

100%|██████████| 10/10 [00:00<00:00, 68.55it/s]  
Train (epoch 63/100) Loss: 0.032019 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 64.51it/s]  
Train (epoch 64/100) Loss: 0.186206 Accuracy: 95.00%  
100%|██████████| 10/10 [00:00<00:00, 71.59it/s]  
Train (epoch 65/100) Loss: 0.086446 Accuracy: 96.00%  
100%|██████████| 10/10 [00:00<00:00, 68.35it/s]  
Train (epoch 66/100) Loss: 0.033519 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 70.66it/s]  
Train (epoch 67/100) Loss: 0.027742 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 71.32it/s]  
Train (epoch 68/100) Loss: 0.078759 Accuracy: 98.00%  
100%|██████████| 10/10 [00:00<00:00, 73.50it/s]  
Train (epoch 69/100) Loss: 0.057207 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 75.39it/s]  
Train (epoch 70/100) Loss: 0.082496 Accuracy: 98.00%  
100%|██████████| 10/10 [00:00<00:00, 68.94it/s]  
Train (epoch 71/100) Loss: 0.070863 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 68.83it/s]  
Train (epoch 72/100) Loss: 0.016160 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 66.12it/s]  
Train (epoch 73/100) Loss: 0.018946 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 70.07it/s]  
Train (epoch 74/100) Loss: 0.037918 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 69.77it/s]  
Train (epoch 75/100) Loss: 0.014793 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 67.54it/s]  
Train (epoch 76/100) Loss: 0.025603 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 61.79it/s]  
Train (epoch 77/100) Loss: 0.047919 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 71.72it/s]  
Train (epoch 78/100) Loss: 0.018200 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 75.68it/s]  
Train (epoch 79/100) Loss: 0.044535 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 72.38it/s]  
Train (epoch 80/100) Loss: 0.077692 Accuracy: 97.00%  
100%|██████████| 10/10 [00:00<00:00, 74.23it/s]  
Train (epoch 81/100) Loss: 0.183003 Accuracy: 94.00%  
100%|██████████| 10/10 [00:00<00:00, 72.80it/s]  
Train (epoch 82/100) Loss: 0.024014 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 74.65it/s]  
Train (epoch 83/100) Loss: 0.033373 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 68.38it/s]  
Train (epoch 84/100) Loss: 0.042100 Accuracy: 97.00%  
100%|██████████| 10/10 [00:00<00:00, 67.83it/s]  
Train (epoch 85/100) Loss: 0.033795 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 72.17it/s]  
Train (epoch 86/100) Loss: 0.014657 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 71.42it/s]  
Train (epoch 87/100) Loss: 0.110096 Accuracy: 98.00%  
100%|██████████| 10/10 [00:00<00:00, 69.25it/s]  
Train (epoch 88/100) Loss: 0.043062 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 69.99it/s]  
Train (epoch 89/100) Loss: 0.047587 Accuracy: 99.00%  
100%|██████████| 10/10 [00:00<00:00, 68.98it/s]  
Train (epoch 90/100) Loss: 0.008716 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 72.40it/s]  
Train (epoch 91/100) Loss: 0.014606 Accuracy: 100.00%  
100%|██████████| 10/10 [00:00<00:00, 72.53it/s]  
Train (epoch 92/100) Loss: 0.028750 Accuracy: 100.00%

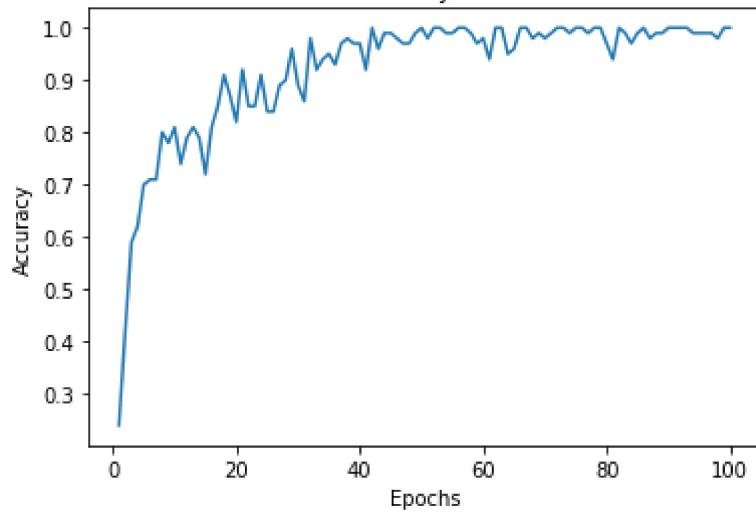
```
100%|██████████| 10/10 [00:00<00:00, 67.11it/s]
Train (epoch 93/100)    Loss: 0.032452 Accuracy: 100.00%
100%|██████████| 10/10 [00:00<00:00, 68.52it/s]
Train (epoch 94/100)    Loss: 0.034519 Accuracy: 99.00%
100%|██████████| 10/10 [00:00<00:00, 65.85it/s]
Train (epoch 95/100)    Loss: 0.047753 Accuracy: 99.00%
100%|██████████| 10/10 [00:00<00:00, 70.55it/s]
Train (epoch 96/100)    Loss: 0.040617 Accuracy: 99.00%
100%|██████████| 10/10 [00:00<00:00, 65.90it/s]
Train (epoch 97/100)    Loss: 0.037906 Accuracy: 99.00%
100%|██████████| 10/10 [00:00<00:00, 65.87it/s]
Train (epoch 98/100)    Loss: 0.129487 Accuracy: 98.00%
100%|██████████| 10/10 [00:00<00:00, 68.00it/s]
Train (epoch 99/100)    Loss: 0.016416 Accuracy: 100.00%
100%|██████████| 10/10 [00:00<00:00, 67.75it/s]
Train (epoch 100/100)   Loss: 0.020519 Accuracy: 100.00%
```

Train loss: 0.020519 Train accuracy: 100.00%

Train loss evolution



Train accuracy evolution



## Testing

In [20]: `test(model, test_dataloader)`

```
100%|██████████| 1000/1000 [00:05<00:00, 189.96it/s]
Test loss: 2.795136 Test accuracy: 29.82%
```

## Results

Model	Number of epochs	Train accuracy	Test accuracy
ResNet18	100	0.64	0.1940
ResNet18 (transfer learning)	100	1.00	0.2982

## Incorporating priors

Geometrical priors are appealing for image classification tasks. A color image  $x$  can be seen as a function:  $\mathbb{S} \rightarrow \mathbb{R}^3$ , where  $\mathbb{S} \subset \mathbb{R}^2$  is the image support. Let us consider transformations  $\mathcal{T}$  of possible inputs  $x$ . For instance, if an image had infinite support, a translation  $\mathcal{T}_a$  of an image  $x$  by a shift  $a$  would lead to a new infinite-support image  $\mathcal{T}_a(x)$ , described at each pixel  $u$  by :

$$\forall u, \mathcal{T}_a(x)(u) = x(u - a).$$

**Question 5 (1.5 points) :** Explain the issues when dealing with translations, rotations, scaling effects, color changes on  $32 \times 32$  images. Propose several ideas to tackle them.

**Answer:**

The problems with transformations such as translations, rotations, scaling effects or color changes on  $32 \times 32$  images is that these transformations introduce some information loss while the original image does not carry much information to begin with. If we use these transformations, the image might be too distorted and the characteristics or the features in the image might not be usable for the model to learn.

To fix this, we could limit ourselves to small translations (or rotations, scaling effects, color changes) to make sure most of the characteristics of the image are preserved. Otherwise, we could increase the image size using some interpolation method.

## Data augmentations

**Question 6 (3 points) :** Propose a set of geometric transformations beyond translation, and incorporate them in your training pipeline. Train the model of the **Question 3** with them and report the accuracies. You can use tools from <https://pytorch.org/vision/stable/transforms.html>

In [11]:

```
train_transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomCrop(32, padding=4),
    transforms.ToTensor(),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
])

test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
])

batch_size = 10

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=train_transform)
trainset = torch.utils.data.Subset(trainset, list(range(100)))
train_dataloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=None)
test_dataloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=True)

len(train_dataloader), len(test_dataloader)
```

Files already downloaded and verified  
 Files already downloaded and verified

Out[11]: (10, 1000)

## Training

In [12]:

```
model = ResNet18()
model.to(device)

epochs = 100
criterion = nn.CrossEntropyLoss()
lr = 0.1
optimizer = optim.Adam(model.parameters(), lr=lr)
lr_update_period = 30
milestones = list(range(0, epochs, lr_update_period))[1:]
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones, gamma=0.1)
```

In [13]:

```
train(model, train_dataloader, criterion, optimizer, scheduler)
```

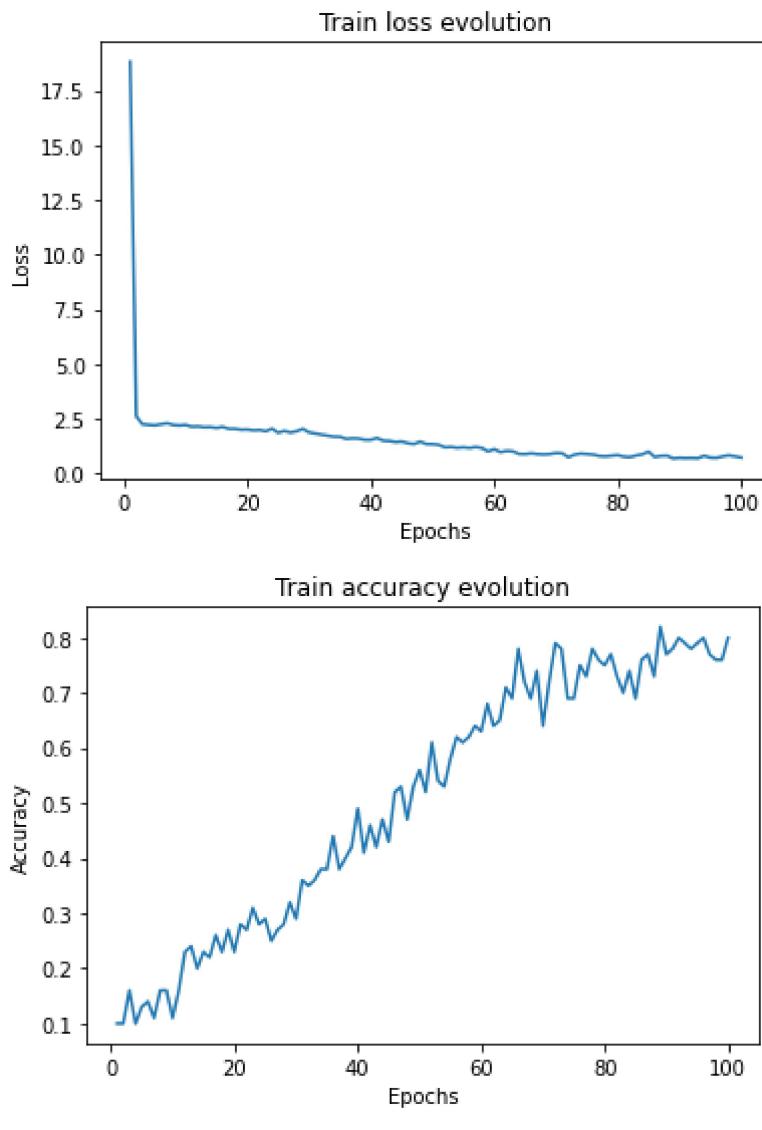
```
100%|██████████| 10/10 [00:00<00:00, 59.71it/s]
Train (epoch 1/100)    Loss: 18.809826 Accuracy: 10.00%
100%|██████████| 10/10 [00:00<00:00, 59.04it/s]
Train (epoch 2/100)    Loss: 2.603609 Accuracy: 10.00%
100%|██████████| 10/10 [00:00<00:00, 58.40it/s]
Train (epoch 3/100)    Loss: 2.242220 Accuracy: 16.00%
100%|██████████| 10/10 [00:00<00:00, 57.67it/s]
Train (epoch 4/100)    Loss: 2.215702 Accuracy: 10.00%
100%|██████████| 10/10 [00:00<00:00, 57.46it/s]
Train (epoch 5/100)    Loss: 2.190791 Accuracy: 13.00%
100%|██████████| 10/10 [00:00<00:00, 57.59it/s]
Train (epoch 6/100)    Loss: 2.240300 Accuracy: 14.00%
100%|██████████| 10/10 [00:00<00:00, 59.23it/s]
Train (epoch 7/100)    Loss: 2.286387 Accuracy: 11.00%
100%|██████████| 10/10 [00:00<00:00, 58.89it/s]
Train (epoch 8/100)    Loss: 2.207250 Accuracy: 16.00%
100%|██████████| 10/10 [00:00<00:00, 57.69it/s]
Train (epoch 9/100)    Loss: 2.186592 Accuracy: 16.00%
100%|██████████| 10/10 [00:00<00:00, 57.58it/s]
Train (epoch 10/100)   Loss: 2.211465 Accuracy: 11.00%
100%|██████████| 10/10 [00:00<00:00, 58.85it/s]
Train (epoch 11/100)   Loss: 2.129247 Accuracy: 16.00%
100%|██████████| 10/10 [00:00<00:00, 59.93it/s]
Train (epoch 12/100)   Loss: 2.140280 Accuracy: 23.00%
100%|██████████| 10/10 [00:00<00:00, 57.62it/s]
Train (epoch 13/100)   Loss: 2.104954 Accuracy: 24.00%
100%|██████████| 10/10 [00:00<00:00, 58.30it/s]
Train (epoch 14/100)   Loss: 2.112455 Accuracy: 20.00%
100%|██████████| 10/10 [00:00<00:00, 59.29it/s]
Train (epoch 15/100)   Loss: 2.072767 Accuracy: 23.00%
100%|██████████| 10/10 [00:00<00:00, 58.22it/s]
Train (epoch 16/100)   Loss: 2.119914 Accuracy: 22.00%
100%|██████████| 10/10 [00:00<00:00, 59.58it/s]
```

Train (epoch 17/100) Loss: 2.032615 Accuracy: 26.00%  
100% [██████████] 10/10 [00:00<00:00, 58.34it/s]  
Train (epoch 18/100) Loss: 2.033696 Accuracy: 23.00%  
100% [██████████] 10/10 [00:00<00:00, 58.02it/s]  
Train (epoch 19/100) Loss: 1.990460 Accuracy: 27.00%  
100% [██████████] 10/10 [00:00<00:00, 59.65it/s]  
Train (epoch 20/100) Loss: 1.999297 Accuracy: 23.00%  
100% [██████████] 10/10 [00:00<00:00, 57.27it/s]  
Train (epoch 21/100) Loss: 1.960430 Accuracy: 28.00%  
100% [██████████] 10/10 [00:00<00:00, 59.66it/s]  
Train (epoch 22/100) Loss: 1.976060 Accuracy: 27.00%  
100% [██████████] 10/10 [00:00<00:00, 58.40it/s]  
Train (epoch 23/100) Loss: 1.919089 Accuracy: 31.00%  
100% [██████████] 10/10 [00:00<00:00, 59.18it/s]  
Train (epoch 24/100) Loss: 2.038855 Accuracy: 28.00%  
100% [██████████] 10/10 [00:00<00:00, 58.85it/s]  
Train (epoch 25/100) Loss: 1.855110 Accuracy: 29.00%  
100% [██████████] 10/10 [00:00<00:00, 58.60it/s]  
Train (epoch 26/100) Loss: 1.934399 Accuracy: 25.00%  
100% [██████████] 10/10 [00:00<00:00, 58.76it/s]  
Train (epoch 27/100) Loss: 1.861685 Accuracy: 27.00%  
100% [██████████] 10/10 [00:00<00:00, 59.62it/s]  
Train (epoch 28/100) Loss: 1.921638 Accuracy: 28.00%  
100% [██████████] 10/10 [00:00<00:00, 59.34it/s]  
Train (epoch 29/100) Loss: 2.029971 Accuracy: 32.00%  
100% [██████████] 10/10 [00:00<00:00, 57.50it/s]  
Train (epoch 30/100) Loss: 1.867102 Accuracy: 29.00%  
100% [██████████] 10/10 [00:00<00:00, 58.41it/s]  
Train (epoch 31/100) Loss: 1.817966 Accuracy: 36.00%  
100% [██████████] 10/10 [00:00<00:00, 59.21it/s]  
Train (epoch 32/100) Loss: 1.764226 Accuracy: 35.00%  
100% [██████████] 10/10 [00:00<00:00, 58.35it/s]  
Train (epoch 33/100) Loss: 1.719914 Accuracy: 36.00%  
100% [██████████] 10/10 [00:00<00:00, 57.06it/s]  
Train (epoch 34/100) Loss: 1.670803 Accuracy: 38.00%  
100% [██████████] 10/10 [00:00<00:00, 58.45it/s]  
Train (epoch 35/100) Loss: 1.668811 Accuracy: 38.00%  
100% [██████████] 10/10 [00:00<00:00, 57.97it/s]  
Train (epoch 36/100) Loss: 1.570626 Accuracy: 44.00%  
100% [██████████] 10/10 [00:00<00:00, 58.85it/s]  
Train (epoch 37/100) Loss: 1.594658 Accuracy: 38.00%  
100% [██████████] 10/10 [00:00<00:00, 56.23it/s]  
Train (epoch 38/100) Loss: 1.585293 Accuracy: 40.00%  
100% [██████████] 10/10 [00:00<00:00, 58.73it/s]  
Train (epoch 39/100) Loss: 1.527768 Accuracy: 42.00%  
100% [██████████] 10/10 [00:00<00:00, 58.28it/s]  
Train (epoch 40/100) Loss: 1.523466 Accuracy: 49.00%  
100% [██████████] 10/10 [00:00<00:00, 57.53it/s]  
Train (epoch 41/100) Loss: 1.611788 Accuracy: 41.00%  
100% [██████████] 10/10 [00:00<00:00, 58.35it/s]  
Train (epoch 42/100) Loss: 1.490047 Accuracy: 46.00%  
100% [██████████] 10/10 [00:00<00:00, 59.19it/s]  
Train (epoch 43/100) Loss: 1.478188 Accuracy: 42.00%  
100% [██████████] 10/10 [00:00<00:00, 57.67it/s]  
Train (epoch 44/100) Loss: 1.427063 Accuracy: 47.00%  
100% [██████████] 10/10 [00:00<00:00, 58.41it/s]  
Train (epoch 45/100) Loss: 1.449704 Accuracy: 43.00%  
100% [██████████] 10/10 [00:00<00:00, 58.24it/s]  
Train (epoch 46/100) Loss: 1.373501 Accuracy: 52.00%

100%|██████████| 10/10 [00:00<00:00, 58.53it/s]  
Train (epoch 47/100) Loss: 1.336140 Accuracy: 53.00%  
100%|██████████| 10/10 [00:00<00:00, 59.37it/s]  
Train (epoch 48/100) Loss: 1.444744 Accuracy: 47.00%  
100%|██████████| 10/10 [00:00<00:00, 56.69it/s]  
Train (epoch 49/100) Loss: 1.336955 Accuracy: 53.00%  
100%|██████████| 10/10 [00:00<00:00, 58.82it/s]  
Train (epoch 50/100) Loss: 1.330273 Accuracy: 56.00%  
100%|██████████| 10/10 [00:00<00:00, 58.61it/s]  
Train (epoch 51/100) Loss: 1.301129 Accuracy: 52.00%  
100%|██████████| 10/10 [00:00<00:00, 58.83it/s]  
Train (epoch 52/100) Loss: 1.183607 Accuracy: 61.00%  
100%|██████████| 10/10 [00:00<00:00, 58.27it/s]  
Train (epoch 53/100) Loss: 1.201599 Accuracy: 54.00%  
100%|██████████| 10/10 [00:00<00:00, 59.23it/s]  
Train (epoch 54/100) Loss: 1.161085 Accuracy: 53.00%  
100%|██████████| 10/10 [00:00<00:00, 57.27it/s]  
Train (epoch 55/100) Loss: 1.185804 Accuracy: 58.00%  
100%|██████████| 10/10 [00:00<00:00, 58.73it/s]  
Train (epoch 56/100) Loss: 1.153504 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 59.23it/s]  
Train (epoch 57/100) Loss: 1.198178 Accuracy: 61.00%  
100%|██████████| 10/10 [00:00<00:00, 59.57it/s]  
Train (epoch 58/100) Loss: 1.155060 Accuracy: 62.00%  
100%|██████████| 10/10 [00:00<00:00, 58.55it/s]  
Train (epoch 59/100) Loss: 0.997102 Accuracy: 64.00%  
100%|██████████| 10/10 [00:00<00:00, 57.66it/s]  
Train (epoch 60/100) Loss: 1.099027 Accuracy: 63.00%  
100%|██████████| 10/10 [00:00<00:00, 59.04it/s]  
Train (epoch 61/100) Loss: 0.966277 Accuracy: 68.00%  
100%|██████████| 10/10 [00:00<00:00, 58.23it/s]  
Train (epoch 62/100) Loss: 1.020356 Accuracy: 64.00%  
100%|██████████| 10/10 [00:00<00:00, 57.71it/s]  
Train (epoch 63/100) Loss: 1.003638 Accuracy: 65.00%  
100%|██████████| 10/10 [00:00<00:00, 58.12it/s]  
Train (epoch 64/100) Loss: 0.883082 Accuracy: 71.00%  
100%|██████████| 10/10 [00:00<00:00, 58.26it/s]  
Train (epoch 65/100) Loss: 0.872688 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 58.46it/s]  
Train (epoch 66/100) Loss: 0.909623 Accuracy: 78.00%  
100%|██████████| 10/10 [00:00<00:00, 58.88it/s]  
Train (epoch 67/100) Loss: 0.873946 Accuracy: 72.00%  
100%|██████████| 10/10 [00:00<00:00, 58.51it/s]  
Train (epoch 68/100) Loss: 0.858894 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 59.25it/s]  
Train (epoch 69/100) Loss: 0.874970 Accuracy: 74.00%  
100%|██████████| 10/10 [00:00<00:00, 58.99it/s]  
Train (epoch 70/100) Loss: 0.923939 Accuracy: 64.00%  
100%|██████████| 10/10 [00:00<00:00, 57.24it/s]  
Train (epoch 71/100) Loss: 0.910861 Accuracy: 72.00%  
100%|██████████| 10/10 [00:00<00:00, 56.16it/s]  
Train (epoch 72/100) Loss: 0.732920 Accuracy: 79.00%  
100%|██████████| 10/10 [00:00<00:00, 57.30it/s]  
Train (epoch 73/100) Loss: 0.846929 Accuracy: 78.00%  
100%|██████████| 10/10 [00:00<00:00, 55.47it/s]  
Train (epoch 74/100) Loss: 0.899868 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 57.36it/s]  
Train (epoch 75/100) Loss: 0.874829 Accuracy: 69.00%  
100%|██████████| 10/10 [00:00<00:00, 58.41it/s]  
Train (epoch 76/100) Loss: 0.852087 Accuracy: 75.00%

```
100%|██████████| 10/10 [00:00<00:00, 58.51it/s]
Train (epoch 77/100)    Loss: 0.789549 Accuracy: 73.00%
100%|██████████| 10/10 [00:00<00:00, 59.51it/s]
Train (epoch 78/100)    Loss: 0.771053 Accuracy: 78.00%
100%|██████████| 10/10 [00:00<00:00, 58.77it/s]
Train (epoch 79/100)    Loss: 0.800555 Accuracy: 76.00%
100%|██████████| 10/10 [00:00<00:00, 58.97it/s]
Train (epoch 80/100)    Loss: 0.826740 Accuracy: 75.00%
100%|██████████| 10/10 [00:00<00:00, 59.17it/s]
Train (epoch 81/100)    Loss: 0.757706 Accuracy: 77.00%
100%|██████████| 10/10 [00:00<00:00, 58.70it/s]
Train (epoch 82/100)    Loss: 0.740442 Accuracy: 73.00%
100%|██████████| 10/10 [00:00<00:00, 56.43it/s]
Train (epoch 83/100)    Loss: 0.809234 Accuracy: 70.00%
100%|██████████| 10/10 [00:00<00:00, 56.66it/s]
Train (epoch 84/100)    Loss: 0.863745 Accuracy: 74.00%
100%|██████████| 10/10 [00:00<00:00, 57.15it/s]
Train (epoch 85/100)    Loss: 0.985828 Accuracy: 69.00%
100%|██████████| 10/10 [00:00<00:00, 58.43it/s]
Train (epoch 86/100)    Loss: 0.745245 Accuracy: 76.00%
100%|██████████| 10/10 [00:00<00:00, 58.31it/s]
Train (epoch 87/100)    Loss: 0.790299 Accuracy: 77.00%
100%|██████████| 10/10 [00:00<00:00, 57.82it/s]
Train (epoch 88/100)    Loss: 0.803151 Accuracy: 73.00%
100%|██████████| 10/10 [00:00<00:00, 57.55it/s]
Train (epoch 89/100)    Loss: 0.675784 Accuracy: 82.00%
100%|██████████| 10/10 [00:00<00:00, 57.34it/s]
Train (epoch 90/100)    Loss: 0.702526 Accuracy: 77.00%
100%|██████████| 10/10 [00:00<00:00, 55.90it/s]
Train (epoch 91/100)    Loss: 0.691293 Accuracy: 78.00%
100%|██████████| 10/10 [00:00<00:00, 58.52it/s]
Train (epoch 92/100)    Loss: 0.697133 Accuracy: 80.00%
100%|██████████| 10/10 [00:00<00:00, 56.09it/s]
Train (epoch 93/100)    Loss: 0.682923 Accuracy: 79.00%
100%|██████████| 10/10 [00:00<00:00, 55.33it/s]
Train (epoch 94/100)    Loss: 0.799321 Accuracy: 78.00%
100%|██████████| 10/10 [00:00<00:00, 57.04it/s]
Train (epoch 95/100)    Loss: 0.713160 Accuracy: 79.00%
100%|██████████| 10/10 [00:00<00:00, 57.86it/s]
Train (epoch 96/100)    Loss: 0.701933 Accuracy: 80.00%
100%|██████████| 10/10 [00:00<00:00, 58.71it/s]
Train (epoch 97/100)    Loss: 0.768586 Accuracy: 77.00%
100%|██████████| 10/10 [00:00<00:00, 55.35it/s]
Train (epoch 98/100)    Loss: 0.819670 Accuracy: 76.00%
100%|██████████| 10/10 [00:00<00:00, 56.23it/s]
Train (epoch 99/100)    Loss: 0.769987 Accuracy: 76.00%
100%|██████████| 10/10 [00:00<00:00, 58.62it/s]
Train (epoch 100/100)   Loss: 0.719964 Accuracy: 80.00%
```

Train loss: 0.719964    Train accuracy: 80.00%



## Testing

```
In [14]: test(model, test_dataloader)
```

```
100%|██████████| 1000/1000 [00:05<00:00, 189.96it/s]
Test loss: 3.474962      Test accuracy: 19.84%
```

## Results

Model	Number of epochs	Train accuracy	Test accuracy
ResNet18	100	0.64	0.1940
ResNet18 (transfer learning)	100	1.00	0.2982
ResNet18 (data augmentation)	100	0.80	0.1984

## Conclusions

**Question 7 (5 points) :** Write a short report explaining the pros and the cons of each method you implemented. 25% of the grade of this project will correspond to this question, thus, it should be done carefully. In particular, please add a plot that will summarize all your numerical results.

**Answer:**

In this notebook, we implemented three training methods on the CIFAR-10 dataset with very few (only 100) training samples while the testing set amounted for 10,000 samples. The accuracies reached by each method are displayed in the following table:

Model	Number of epochs	Train accuracy	Test accuracy
ResNet18	100	0.64	0.1940
ResNet18 (transfer learning)	100	1.00	0.2982
ResNet18 (data augmentation)	100	0.80	0.1984

## ResNet18 model trained from scratch

The first training method we implemented was training the ResNet18 model from scratch (i.e. from randomly initialized weights). This model reached a training accuracy of 64% after 100 epoch but its accuracy on the testing set was only 19.40%.

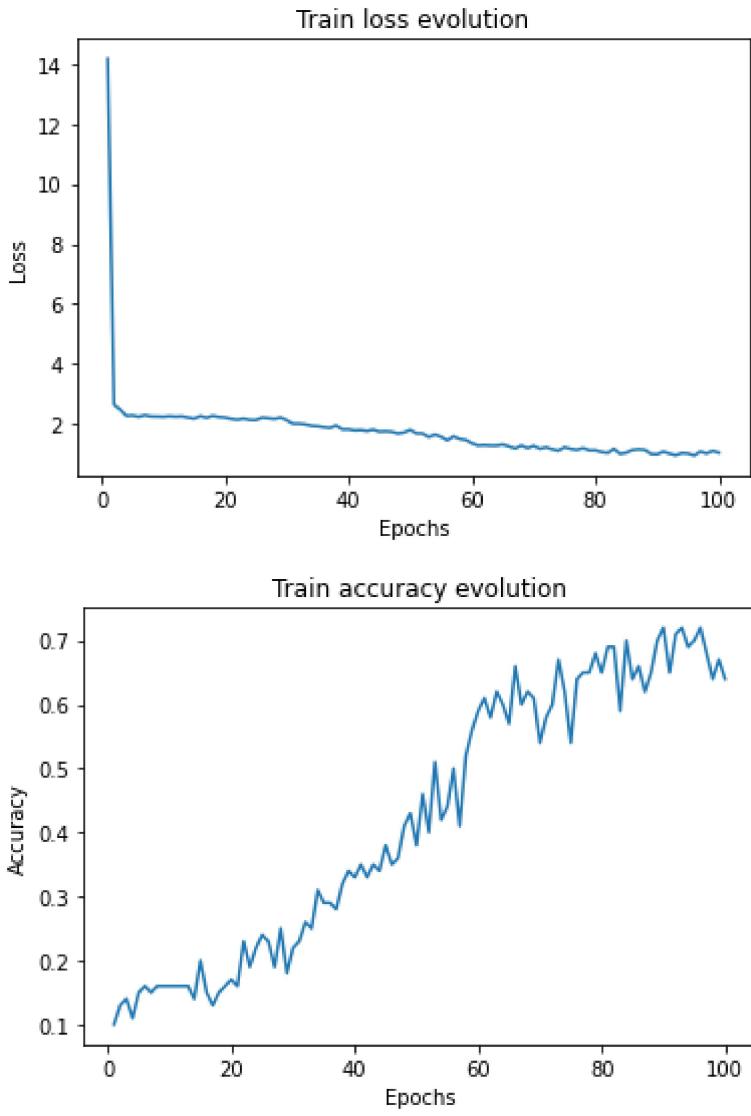
The advantages of this method are the following:

- Does not use pre-trained weights so the model architecture is free to be outside of the typically trained models.

The drawbacks of this method are the following:

- Easily overfits the training set, especially if there are few training samples. Here the training loss was 1.043173 in the end while the testing loss was 2.650611 (i.e. around 2.5 times higher) so we can guess that the overfitting started way before the 100th epoch;
- As shown in the accuracy graph bellow, this method takes a longer time (high number of epochs) to reach a good accuracy. After 60 epochs, the loss is already quite lower than at the start but the accuracy is only around 50%.

Training loss and accuracy graphs:



## ResNet18 model trained by transfer learning

The second training method we implemented was training the same model (ResNet18) by transfer learning using weights pre-trained on the ImageNet dataset. We choose to use this same model to be able to reliably compare the results of our experiments. With this method, the training accuracy after 100 epochs was 100% while the testing accuracy was 29.82%.

The advantages of this method are the following:

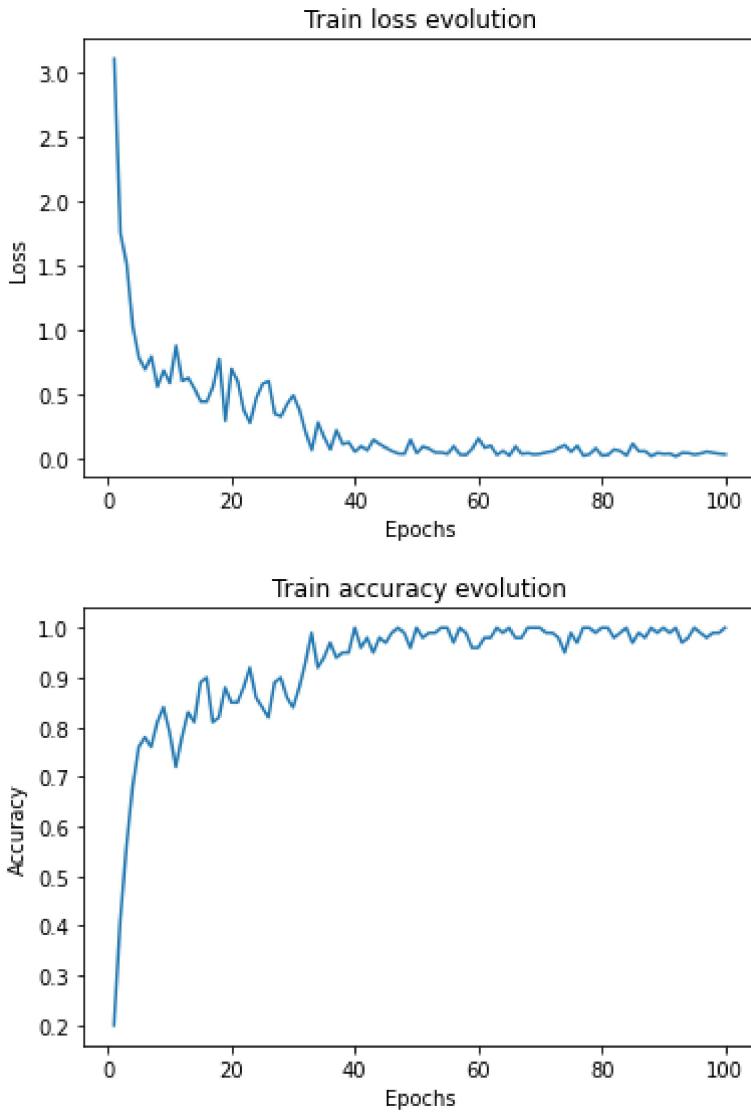
- It reaches a good training accuracy after only a few epoch. As shown in the accuracy graph bellow, the model reaches a training accuracy over 80% after only 20 epochs and at the 40th epoch it reached 100% accuracy already;
- It allows the model to reach lower loss since the basic image features are already known by the pre-trained model. The training loss ended up at 0.020519 which is 5 times lower than with the previous method.

The drawbacks of this method are the following:

- This method is still subject to overfitting. We can see it in the clear gap in training and testing accuracies but also in the fact that the model reaches 100% training accuracy after 40 epochs. We need to carefully fine tune the learning rate and number of epochs in order to reduce this risk;

- The architecture of the model in use has to be part of the typical models trained on ImageNet (or any other dataset). Many pre-trained models are available but it does not allow for unlimited freedom in the architecture of the model.

Training loss and accuracy graphs:



## ResNet18 model trained from scratch using data augmentation

The third and last training method we implemented was adding a data augmentation step in the first training method. We used an horizontal flip (no vertical flip to avoid upside down images), a random crop and a normalization of the pixel colors in a preprocessing pipeline. This method reached similar results to the first method with a training accuracy of 80% and a testing accuracy of 19.84%.

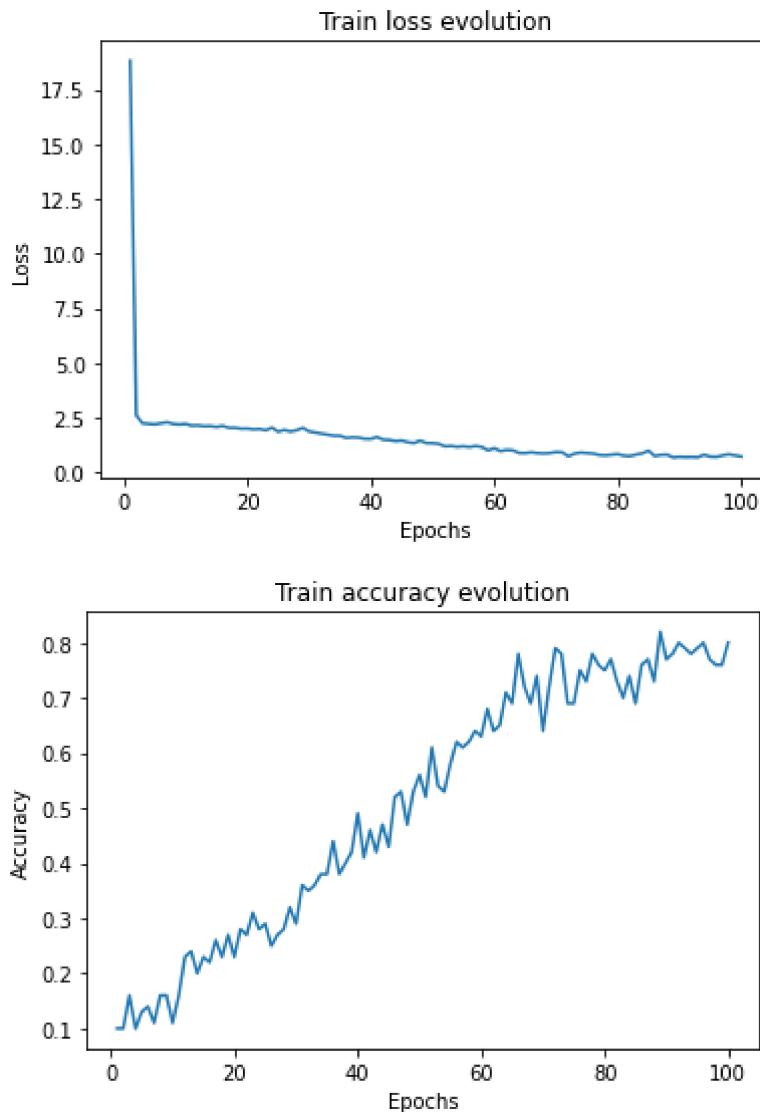
The advantages of this method are the following:

- It can improve the performance of learning processes on small datasets by generating new labeled samples;
- Does not use pre-trained weights so the model architecture is free to be outside of the typically trained models.

The drawbacks of this method are the following:

- With small images such as the ones used here, many transformations are dangerous to use such as translations, rotations, scaling effects or color changes. Thus, the benefits of this method are limited here;
- This method is also subject to overfitting. Here the model reached better training accuracy and loss (0.719964) than the first method but similar testing accuracy and poorer loss (3.474962). This shows that the data augmentation was not successful with only 100 samples of  $32 \times 32$  images.

Training loss and accuracy graphs:



## Weak supervision

**Bonus [open] question (up to 4 points)** : Pick a weakly supervised method that will potentially use  $\mathcal{X}_{\text{nolabel}} \cup \mathcal{X}_{\text{train}}$  to train a model (a subset of  $\mathcal{X}_{\text{nolabel}}$  is also fine). Evaluate it and report the accuracies. You should be careful in the choice of your method, in order to avoid heavy computational effort.